

Assignment-6: CNN

ID: 2023E8014682040, 邓勇, 自动化研究所

Modules of CNN

下面是关于卷积神经网络的相关内容:

- Convolutional Layer
 - input size: $(n_h \times n_w)$, convo kernel: $(k_h \times k_w)$,
 - padding: (p_h, p_w) , stride: (s_h, s_w)
 - output size:

$$(n_h - k_h + p_h + 1) \times (n_w - k_w + p_w + 1)$$

- output size:

$$\lfloor (n_h - k_h + p_h + s_h) / s_h \rfloor \times \lfloor (n_w - k_w + p_w + s_w) / s_w \rfloor$$

, or we can just compute

$$\lfloor (N - K + P) / S \rfloor + 1$$

if $p_h = k_h - 1, p_w = k_w - 1$, output $(n_h / s_h) \times (n_w / s_w)$

- Max Pooling
 - (2 x 2) max pooling, output size: $(n_h / 2) \times (n_w / 2)$
- Pytorch code
 - torch.nn.Conv2d : 卷积层
 - torch.nn.MaxPool2d : 最大池化层
 - torch.nn.Flatten : 展平层
 - torch.nn.Linear : 全连接层
 - torch.nn.Sequential : 顺序模型
- utils
 - out = torchvision.utils.make_grid(images) : 显示图像
 - imshow(out, title=[class_names[x] for x in classes]) : 显示图像
 - torchvision.transforms.ToPILImage() : 图像转换

Designed convolutional layer as following:

```
class ConvNetwork(nn.Module):
    # original shape of images [bz, 3, 28, 28]
    # input_layer: 3 input channels, 6 output channels, 5 kernel size
    def __init__(self):
        super(ConvNetwork, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, 5) # [bz, 3, 28, 28] -> [bz, 6, 24, 24]
        self.pool = nn.MaxPool2d(2, 2) # [bz, 6, 24, 24] -> [bz, 6, 12, 12]
        self.conv2 = nn.Conv2d(6, 16, 3) # [bz, 6, 12, 12] -> [bz, 16, 10, 10]
        # pool: [bz, 16, 10, 10] -> [bz, 16, 5, 5]
        self.fc1 = nn.Linear(16 * 5 * 5, 120) # [bz, 16, 5, 5] -> [bz, 120]
        self.fc2 = nn.Linear(120, 84) # [bz, 120] -> [bz, 84]
        self.fc3 = nn.Linear(84, 10) # [bz, 84] -> [bz, 10]

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool(x)
        x = F.relu(self.conv2(x))
        x = self.pool(x)
        # flatten the output of conv2 to (batch_size, 16*5*5)
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        return self.fc3(x)
```

Training and Testing using Pytorch

```
writer = SummaryWriter("runs/mnist")

# device config
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# Hyperparameters
input_size = 784 # 28x28
hidden_size = 50
num_classes = 10
num_epochs = 5
batch_size = 32
learning_rate = 0.001

# MNIST dataset
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, ), (0.5, ))]
)

train_dataset = datasets.MNIST(root='./data', train=True, transform=transform, download=True)
test_dataset = datasets.MNIST(root='./data', train=False, transform=transform, download=True)
train_loader = DataLoader(dataset=train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(dataset=test_dataset, batch_size=batch_size, shuffle=False)

examples = iter(train_loader)
images, labels = examples.__next__()
print(f'Size of data, labels: {images.shape}, {labels.shape}')

img_grid = torchvision.utils.make_grid(images)
writer.add_image('MNIST Images', img_grid)
```

我们可以看到它读入图片和标签：




fig1: Batch data of MNIST

接下来我们训练和测试该模型：

```

model = ConvNetwork()
# loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

# training loop
n_total_steps = len(train_loader)
running_loss = 0.0
running_correct = 0

for epoch in range(num_epochs):
    for i, (images, labels) in enumerate(train_loader):
        images = images.to(device)
        labels = labels.to(device)

        # forward pass
        outputs = model(images)
        loss = criterion(outputs, labels)
        # backward pass and update weights
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # tensorboard
        running_loss += loss.item()
        _, predicted = torch.max(outputs.data, 1)
        running_correct += (predicted == labels).sum().item()

    if (i+1) % 100 == 0:
        print(f'epoch {epoch+1}/{num_epochs}, step {i+1}/{len(train_loader)}, loss = {running_l
writer.add_scalar('training loss', running_loss/100, epoch*n_total_steps + 1)
writer.add_scalar('accuracy', running_correct/100, epoch*n_total_steps + 1)
        running_loss = 0.0
        running_correct = 0

# Test
preds = []
pred_labels = []

with torch.no_grad():
    n_correct = 0
    n_samples = 0
    for images, labels in test_loader:
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)

        # max returns (value, index)

```

```
_, predicted = torch.max(outputs.data, 1)
n_samples += labels.shape[0]
n_correct += (predicted == labels).sum().item()

# classification results for tensorboard
class_predictions = [F.softmax(output, dim=0) for output in outputs]
# print(predicted, labels)
preds.append(class_predictions)
pred_labels.append(predicted)

preds = torch.cat([torch.stack(batch) for batch in preds])
pred_labels = torch.cat(pred_labels, dim=0)
acc = 100.0 * n_correct / n_samples
print(f'Accuracy on the testing images = {acc}%',)

for i in range(num_classes):
    labels_i = pred_labels == i
    preds_i = preds[:, i]
    writer.add_pr_curve(str(i), labels_i, preds_i, global_step=0)

writer.close()
```

Pytorch Lightning Implementation

PyTorch Lightning 是一个用于深度学习项目的轻量级开发框架，它构建在 PyTorch 之上，提供了更高级别的抽象。PyTorch Lightning 的设计理念是尽可能地减少样板代码，让用户专注于模型的设计和调试，同时提供了丰富的功能和扩展接口，使得用户可以根据自己的需求定制训练过程。下面通过 Pytorch Lightning 实现一个简单的 CNN 模型。

```

import torch
import torch.nn as nn
import torch.nn.functional as F
import pytorch_lightning as pl
from torchvision import transforms, datasets
from torch.utils.data import DataLoader, Dataset

transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, ), (0.5, ))]
)

# Hyperparameters

num_epochs = 5
batch_size = 100
learning_rate = 0.001

class LitCNN(pl.LightningModule):
    # original shape of images [bz, 1, 28, 28]
    # input_layer: 3 input channels, 6 output channels, 5 kernel size
    def __init__(self):
        super(LitCNN, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, 5) # [bz,1, 28, 28] -> [bz, 6, 24, 24]
        self.pool = nn.MaxPool2d(2, 2) # [bz, 6, 24, 24] -> [bz, 6, 12, 12]
        self.conv2 = nn.Conv2d(6, 16, 3) # [bz, 6, 12, 12] -> [bz, 16, 10, 10]
        # pool: [bz, 16, 10, 10] -> [bz, 16, 5, 5]
        self.fc1 = nn.Linear(16 * 5 * 5, 120) # [bz, 16, 5, 5] -> [bz, 120]
        self.fc2 = nn.Linear(120, 84) # [bz, 120] -> [bz, 84]
        self.fc3 = nn.Linear(84, 10) # [bz, 84] -> [bz, 10]

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool(x)
        x = F.relu(self.conv2(x))
        x = self.pool(x)
        # flatten the output of conv2 to (batch_size, 16*5*5)
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        return self.fc3(x)

    def configure_optimizers(self):
        return torch.optim.SGD(self.parameters(), lr=learning_rate)

    def training_step(self, batch, batch_idx):
        images, labels = batch

```

```
        outputs = self(images)
        loss = F.cross_entropy(outputs, labels)
        tensorboard_logs = {'train_loss': loss}
        return {'loss': loss, 'log': tensorboard_logs}

def train_dataloader(self):
    train_dataset = datasets.MNIST(root='./data', train=True, transform=transform, download=True)
    train_loader = DataLoader(dataset=train_dataset, batch_size=batch_size, shuffle=True)
    return train_loader

model = LitCNN()
trainer = pl.Trainer(max_epochs=num_epochs, fast_dev_run=False)
trainer.fit(model)
```

Results

训练超参数: num_epochs=20, batchsize=200
CNN 训练结果




fig2: CNN training loss and accuracy

CNN 测试结果:
Accuracy on the testing images = 43.94%